

# OpenMP

## An Easy Way for Parallel Programming

Chuizheng Meng

May 17, 2016

# What is OpenMP

- ▶ OpenMP is a specification for a set of compiler directives, library routines, and environment variables that can be used to specify high-level parallelism in Fortran and C/C++ programs.
- ▶ OpenMP is designed for Fortran, C and C++.
- ▶ Supported compilers
  - ▶ GCC( $\geq 4.2.0$ )
  - ▶ Visual Studio 2008-2010 C++
  - ▶ Clang( $\geq 3.8$ )
  - ▶ IBM, Intel, Texas Instrument...
- ▶ Only a supported compiler is needed!

# Hello OpenMP

## hello.cpp

```
1  #include <stdio.h>
2
3  int main() {
4      #pragma omp parallel
5      printf("Hello OpenMP!\n");
6      return 0;
7  }
```

## Compile & Run(with flag "-fopenmp")

```
1  g++ -Wall -Werror -O3 hello.cpp -o hello -fopenmp
2  ./hello
```


# Hello OpenMP

## Result

```
neozero ~ > oop2016 > week13pre > ./hello
Hello OpenMP!
Hello OpenMP!
Hello OpenMP!
Hello OpenMP!
neozero ~ > oop2016 > week13pre > 
```

# How OpenMP Works

## Fork-Join Model



- ▶ All OpenMP programs begin as a single process: the master thread.
- ▶ Fork
- ▶ Join

# How OpenMP Works

- ▶ Shared Memory Model
- ▶ Thread Based Parallelism
- ▶ Explicit Parallelism
- ▶ Compiler Directive Based
- ▶ Dynamic Threads
- ▶ I/O(It is your duty!)

# OpenMP API Overview

- ▶ Compiler Directives
- ▶ Runtime Library Routines
- ▶ Environment Variables

# Compiler Directives

## Format

`#pragma omp directive-name [clause,...]`



# Directive-name

- ▶ parallel
  - ▶ A parallel region is a block of code that will be executed by multiple threads.
  - ▶ It is the fundamental OpenMP parallel construct.
- ▶ for
  - ▶ Shares iterations of a loop across the team.
  - ▶ Data parallelism
- ▶ sections
  - ▶ Each section is executed by a thread.
  - ▶ Functional parallelism
- ▶ single
  - ▶ The enclosed code is to be executed by only one thread in the team. Other threads will never execute it.
  - ▶ Can be used to deal with IO.

# Directive-name

- ▶ critical
  - ▶ Only ONE thread can execute codes in the block at same time.
  - ▶ Often used when a shared variable is modified.
- ▶ atomic
  - ▶ A unit of storage can only be modified by ONE thread at same time.
  - ▶ Often used when some value of a shared array is modified.

# Clause

Clause	Directive					
	PARALLEL	DO/for	SECTIONS	SINGLE	PARALLEL DO/for	PARALLEL SECTIONS
IF	●				●	●
PRIVATE	●	●	●	●	●	●
SHARED	●	●			●	●
DEFAULT	●				●	●
FIRSTPRIVATE	●	●	●	●	●	●
LASTPRIVATE		●	●		●	●
REDUCTION	●	●	●		●	●
COPYIN	●				●	●
COPYPRIVATE				●		
SCHEDULE		●			●	
ORDERED		●			●	
NOWAIT		●	●	●		

# Clause




- ▶ `if (expression)`
  - ▶ True: A team of threads is created.
  - ▶ False: The region is executed serially by the master thread.
- ▶ `private(list)`
  - ▶ Variables in the list is private to each thread.
  - ▶ They should be assumed to be uninitialized for each thread.(Use `firstprivate` to solve that.)
- ▶ `shared(list)`
  - ▶ Variables in the list is shared among all threads.
  - ▶ It is the programmer's responsibility to ensure that multiple threads properly access shared variables!

# Clause

- ▶ `reduction(op:list)`
  - ▶ Performs a reduction on the variables that appear in its list.
  - ▶ Operator must observe commutative law and associative law.
  - ▶ For some variable  $x$  in list, multiple private copies of  $x$  is created in each thread, and finally  $x$  is set as  $(x \text{ op } x \text{ op } x \text{ op } \dots \text{ op } x)$ .
- ▶ `schedule(kind[, chunksize])`
  - ▶ Automatically arrange "working schedule" for different threads of for loop.
  - ▶ Chunksize: size of each piece of work.
  - ▶ Kind: method to distribute work, including static, dynamic, guided and runtime.

# Clause

- ▶ `schedule(kind[, chunksize])`
  - ▶ Kind
    - ▶ static: loop
    - ▶ dynamic: first come, first serve
    - ▶ guided: like dynamic, but with descending chunksize
    - ▶ runtime: follow the env value `OMP_SCHEDULE`



# Clause

## Static

```
neozero ~ > oop2016 > week13pre > ./reporter
我是香港记者, +4s
我是香港记者, +5s
我是香港记者, +6s
我是西方记者, +7s
我是西方记者, +8s
我是西方记者, +9s
我是大陆记者, +1s
我是大陆记者, +2s
我是大陆记者, +3s
我是华莱士, +10s
我是华莱士, +11s
我是华莱士, +12s
```

# Clause

## Dynamic

```
neozero ~ > oop2016 > week13pre > ./reporter  
我是香港记者, +1s  
我是香港记者, +5s  
我是香港记者, +6s  
我是香港记者, +7s  
我是香港记者, +8s  
我是香港记者, +9s  
我是香港记者, +10s  
我是香港记者, +11s  
我是香港记者, +12s  
我是华莱士, +4s  
我是西方记者, +3s  
我是大陆记者, +2s
```



# Clause

## Guided

```
neozero ~ > oop2016 > week13pre > ./reporter
我是大陆记者, +1s
我是大陆记者, +2s
我是香港记者, +7s
我是香港记者, +8s
我是香港记者, +10s
我是大陆记者, +3s
我是大陆记者, +12s
我是西方记者, +4s
我是西方记者, +5s
我是西方记者, +6s
我是香港记者, +11s
我是华莱士, +9s
```

# Runtime Library Routines

- ▶ For C/C++, include the `<omp.h>` header file.
- ▶ `void omp_set_num_threads(int num_threads)`
- ▶ `int omp_get_num_threads(void)`
- ▶ `int omp_get_thread_num(void)`
- ▶ Manipulations about lock(using critical may be easier)
- ▶ ...

# Environment Variables

- ▶ OpenMP provides some environment variables for controlling the execution of parallel code.
- ▶ OMP\_SCHEDULE
- ▶ OMP\_NUM\_THREADS
- ▶ OMP\_THREAD\_LIMIT

# Examples

- ▶ Merge Sort with OpenMP
- ▶ Numeric calculation of  $\pi$

# Merge Sort

```
void mergesort(int* array, int start, int end, int
    thread, int* buffer) {
    if (end - start <= 1)
        return;
    int p = ((start + end) >> 1);
    if (thread <= 1) {
        mergesort(array, start, p, 1, buffer);
        mergesort(array, p, end, 1, buffer);
    }
```

# Merge Sort

```
else {
    #pragma omp parallel sections
    {
        #pragma omp section
        {
            mergesort(array, start, p, thread / 2, buffer);
        }
        #pragma omp section
        {
            mergesort(array, p, end, thread - thread / 2,
                      buffer);
        }
    }
    merge(array, start, p, end, buffer);
}
```

# Merge Sort

```
neozero ~ > oop2016 > week13pre > gcc -O3 mergesort_demo.c -o mergesort_demo  
-fopenmp  
neozero ~ > oop2016 > week13pre > gcc -O3 mergesort_demo.c -o mergesort_demo_  
nomp  
neozero ~ > oop2016 > week13pre > ./mergesort_demo 100000000 master  
Time cost: 8.189 s  
neozero ~ > oop2016 > week13pre > ./mergesort_demo_nomp 100000000  
Time cost: 14.674 s  
neozero ~ > oop2016 > week13pre > [ ] master
```

# Numeric calculation of $\pi$

$$\int_0^1 \frac{4}{1+x^2} = 4 \arctan 1 = \pi \quad (1)$$

$$\int_0^1 \frac{4}{1+x^2} = \lim_{N \rightarrow +\infty} \frac{1}{N} \sum_{i=0}^{N-1} \frac{4}{1 + \left(\frac{i}{N}\right)^2} \quad (2)$$



# Numeric calculation of $\pi$

```
#pragma omp parallel for schedule(static) reduction(+:  
    sum)  
for (i = 0; i < MAXSTEP; ++i) {  
    sum = sum + FUNC((double)i/MAXSTEP);  
}
```

```
neozero ~ > oop2016 > week13pre > gcc -O3 pi_calc.c -o pi_calc -fopenmp  
neozero ~ > oop2016 > week13pre > gcc -O3 pi_calc.c -o pi_calc_withoutmp  
neozero ~ > oop2016 > week13pre > ./pi_calc  
Calculated PI is: 3.1415926546  
Time: 7.164 s  
neozero ~ > oop2016 > week13pre > ./pi_calc_withoutmp  
Calculated PI is: 3.1415926546  
Time: 13.500 s  
neozero ~ > oop2016 > week13pre >
```

# Pros & Cons

## ► Pros

- Make full use of CPU with little work.
- Portable multithreading code, not platform-specific.
- Work can be easily scheduled.
- Unified code for both serial and parallel applications.

## ► Cons

- High chance of accidentally writing false sharing code.
- Difficult to debug.
- Not very suitable for some situation, i.e. network programming.
- Mainly for parallel computing, not for distributed computing.

# Reference

- ▶ "OpenMP" - Blaise Barney, Lawrence Livermore National Laboratory  
<https://computing.llnl.gov/tutorials/openMP/>
- ▶ OpenMP 并行编程 - 中国科学技术大学超级计算中心  
<http://scc.ustc.edu.cn/zlsc/cxyy/200910/W020121113517997951933.pdf>
- ▶ OpenMP - Wikipedia  
<https://en.wikipedia.org/wiki/OpenMP>